

Using neural networks for classification and regression problems

Eina B. Jørgensen, Anna Lina P. Sjur and Jan-Adrian H. Kallmyr

November 7, 2019

Abstract

1 Introduction

2 Theory

2.1 Logistic Regression

Classification problems aim to predict the behaviour of a given object, and look for patterns based on discrete variables (i.e categories). Logistic regression can be used to solve such problems, commonly by the use of variables with binary outcomes such as true/false, positive/negative, success/failure etc., or in the specific credit card case: *risky/non-risky*

As opposed to linear regression, the equation one gets as a result of minimization of the cost function by $\hat{\beta}$ using logistic regression, is non-linear, and is solved using minimization algorithms called *gradient descent methods*.

When predicting the the output classes in which an object belongs, the prediction is based on the design matrix $\hat{\mathbf{X}} \in \mathbb{R}^{n \times p}$ that contain n samples that each carry p features.

A distinction is made between *hard classification* - deterministically determine the variable to a category, and *soft classification* - determines the probability that a given variable belongs in a certain category. The latter is favorable in many cases, and logistic regression is the most used example of this type of classifier.

When using logistic regression, the probability that a given data point x_i belongs in a category y_i is given by the Sigmoid-function (or logistic function):

$$p(t) = \frac{1}{1 + e^{-t}} = \frac{e^t}{1 + e^t} \quad (1)$$
$$1 - p(t) = p(-t)$$

Assuming a binary classification problem, i.e. y_i can be either 0 or 1, and a set of predictors $\hat{\beta}$ the Sigmoid function (1) gives the probabilities with relation:

$$p(y_i = 0|x_i, \hat{\beta}) = 1 - p(y_i = 1|x_i, \hat{\beta})$$

The total likelihood for all possible outcomes $\mathcal{D} = \{(y_i, x_i)\}$ is used in the Maximum Likelihood Estimation (MLE), aiming at maximizing the log/likelihood function (2). The likelihood function can be expressed with \mathcal{D} :

$$P(\mathcal{D}|\hat{\beta}) = \prod_{i=1}^n \left[p(y_i = 1|x_i, \hat{\beta}) \right]^{y_i} \left[1 - p(y_i = 0|x_i, \hat{\beta}) \right]^{1-y_i}$$

And the log/likelihood function is then:

$$P_{\log}(\hat{\beta}) = \sum_{i=1}^n \left(y_i \log \left[p(y_i = 1|x_i, \hat{\beta}) \right] + (1 - y_i) \log \left[1 - p(y_i = 0|x_i, \hat{\beta}) \right] \right) \quad (2)$$

The cost/error-function \mathcal{C} (also called cross-entropy in statistics) is the negative of the log/likelihood. Maximizing P_{\log} is thus the same as minimizing the cost function. The cost function is:

$$\begin{aligned} \mathcal{C}(\hat{\beta}) = -P_{\log}(\hat{\beta}) = \\ - \sum_{i=1}^n \left(y_i \log [p(y_i = 1|x_i, \hat{\beta})] \right. \\ \left. + (1 - y_i) \log [1 - p(y_i = 0|x_i, \hat{\beta})] \right) \end{aligned} \quad (3)$$

Finding the parameters $\hat{\beta}$ that minimize the cost function is then done through derivation. Defining the vector \hat{y} containing n elements y_i , the $n \times p$ matrix \hat{X} containing the x_i elements, and the vector \hat{p} that is the fitted probabilities $p(y_i|x_i, \hat{\beta})$, the first derivative of \mathcal{C} is

$$\frac{\partial \mathcal{C}(\hat{\beta})}{\partial \hat{\beta}} = -\hat{X}^T(\hat{y} - \hat{p}) \quad (4)$$

This gives rise to set of linear equations, where the aim is to solve the system for $\hat{\beta}$. By introduction of a diagonal matrix \hat{W} with diagonal elements $p(y_i|x_i, \hat{\beta}) \cdot (1 - p(y_i|x_i, \hat{\beta}))$ the second derivative is:

$$\frac{\partial^2 \mathcal{C}(\hat{\beta})}{\partial \hat{\beta} \partial \hat{\beta}^T} = \hat{X}^T \hat{W} \hat{X} \quad (5)$$

With $\hat{x} = [1, x_1, x_2, \dots, x_p]$ and p predictors $\hat{\beta} = [\beta_0, \beta_1, \beta_2, \dots, \beta_p]$ the ration between likelihoods of outcome is:

$$\log \frac{p(\hat{\beta}\hat{x})}{1 - p(\hat{\beta}\hat{x})} = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p \quad (6)$$

and $p(\hat{\beta}\hat{x})$ defined by:

$$p(\hat{\beta}\hat{x}) = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p}} \quad (7)$$

2.2 Gradient Descent Methods

2.3 Neural networks

In this section, the equations used are based off the book by [Nielsen \(2015\)](#).

The structure of a network

Neural networks, as the name suggests, are inspired by our understanding of how networks of neurons function in the brain. As can be seen in the example network in Figure 1, neurons are structured in layers. We always have a input and an output layer, in addition to a varying number of hidden layers. The input layer has as many neurons as there are input variables, while the output layer has one neuron for each output. How many neurons you have in the output layer depends on the specific problem. The number of neurons in each hidden layer, on the other hand, is not directly related to inputs or outputs, and must be decided in some other way.

As the diagram in Figure 1 suggests, the neurons in each layer are not connected with each other, but takes in inputs from the previous layer and passes on an output to the neurons in the next layer, as illustrated with arrows. This way, the inputs are fed through the network and processed, resulting in an output.




Figure 1: Schematic diagram of a neural network with five input neurons in the input layer, one hidden layer with three neurons and a single output neuron in the output layer.

Forward feeding

Each neuron has one or multiple inputs, as illustrated with arrows in Figure 1. Each of these inputs has a weight associated with it. To clarify the notation used, let's take a look at the j th neuron in the l th layer. The weight associated with the input coming from the k th neuron in the previous layer is denoted as w_{jk}^l . In addition, each neuron has a bias associated with it, for the neuron in question denoted as b_j^l . Summing the weighted inputs and the bias, and feeding this to a function σ , gives the activation a_j^l :

$$a_j^l = \sigma \left(\left(\sum_k w_{jk}^l a_k^{l-1} \right) + b_j^l \right)$$

This activation is then fed forward as input to all the neuron in the next layer.

In matrix notation, the activation for the whole layer l can be written as

$$\mathbf{a}^l = \sigma \left(\mathbf{w}^l \mathbf{a}^{l-1} + \mathbf{b}^l \right) \quad (8)$$

Here, \mathbf{a}^l and \mathbf{b}^l are vertical vectors containing the activations and biases of the l th layer, while \mathbf{w}^l is a matrix with elements w_{jk}^l , i.e. the j th column contains the weights of the inputs reaching the j th neuron.

Let's look at the activation function in Eq. (8) denoted with a σ . The use of σ as notation is not arbitrary, since the sigmoid function stated in Eq. (1) is often used. As we will see in the backpropagation algorithm, the sigmoid is a good choice for activation function, since a small change in the output can be propagated backwards, resulting in small changes in the weights and biases through the network.

With a basis in Eq. (8), the algorithm for forward feeding is given in Algorithm 1. Here L is the total number of layers.

```
Set  $\mathbf{a}^1 = \text{input}$ ;  
foreach  $l=2:L$  do  
  | Compute  $\mathbf{a}^l$ ;  
Set output to  $\mathbf{a}^L$ ;
```

Algorithm 1: The forward feeding algorithm.

Backpropagation

When training the network, the goal is to find the weights and biases that minimize the cost function C . For a classification problem, the cost function is often given as

$$C = \frac{1}{2} \sum_i (a_i^L - y_i)^2 \quad (9)$$

To minimize Eq.(9), one can use Stochastic Gradient Decent, as described previously. But in order to use SGD, the derivatives of C must be computed, and it is here that backpropagation comes in. It can be shown that the derivatives are given as in Eq. (10). For a derivation of these expressions see APPENDIX?!?!?!

$$\begin{aligned} \delta^L &= \nabla_a C \odot \sigma'(z^L) \delta^L \\ \delta^l &= ((\mathbf{w}^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l) \\ \frac{\partial C}{\partial b_j^l} &= \delta_j^l \\ \frac{\partial C}{\partial w_{jk}^l} &= a_k^{l-1} \delta_j^l \end{aligned} \quad (10)$$

Adapting neural networks to regression

In order to adapt the network to regression, some changes must be made in Algorithm 1 and 2.

```

Compute  $\{\mathbf{a}^l\}_{l=1}^L$  with feed forward;
Compute  $\delta^L$ ;
Set  $\frac{\partial C}{\partial \mathbf{b}^L} = \delta^L$ ;
Compute  $\frac{\partial C}{\partial \mathbf{w}^L} = \delta^L (\mathbf{a}^{L-1})^T$ ;
foreach  $l=L-1:2$  do
    Compute  $\delta^l$ ;
    Set  $\frac{\partial C}{\partial \mathbf{b}^l} = \delta^l$ ;
    Compute  $\frac{\partial C}{\partial \mathbf{w}^l} = \delta^l (\mathbf{a}^{l-1})^T$ ;

```

Algorithm 2: The backpropagation algorithm.

Figure 2

Implementation

3 Results

4 Discussion

5 Conclusion

References

Michael A. Nielsen. *Neural Networks and Deep Learning*. Determination Press, 2015.